

Fraudulent Claim Detection – Project Report

Problem Statement

Insurance fraud leads to significant financial losses and operational inefficiencies for insurance companies. Detecting fraudulent claims at an early stage is critical to reduce payouts on illegitimate claims while ensuring genuine claims are processed smoothly.

The objective of this assignment is to build and evaluate machine learning models to predict whether an insurance claim is fraudulent or not, based on customer, policy, incident, and claim-related attributes. The focus is not only on predictive accuracy, but also on evaluating business-relevant metrics such as recall (sensitivity), which is particularly important in fraud detection scenarios.

Overall Approach

The solution follows a structured machine learning workflow:

1. Data Load & Preparation
2. Data Cleaning
3. Train-Validation Split
4. EDA on Training Data
5. EDA on Validation Data
6. Feature Engineering
7. Model Building (Logistic, Random Forest & Hyper-parameter Tuning)
8. Prediction & Model Evaluation
9. Conclusion & Key Insights

Each step is implemented incrementally and evaluated using appropriate statistical and classification metrics.

Detailed Explanation

1. Data Load & Preparation

The insurance claims dataset was loaded into the notebook using pandas. Initial inspection included:

- Viewing the first few rows (`df.head()`)
- Checking dataset dimensions
- Inspecting column names and data types

The dataset initially contained **1000 rows and 40 columns**, including customer demographics, policy details, incident attributes, and claim amount components. The target variable `fraud_reported` indicated whether a claim was fraudulent (Y) or not (N).

2. Data Cleaning

2.1 Missing Value Analysis

Missing value checks were performed using:

```
df.isnull().sum()
```

This analysis confirmed that **all columns except _c39 had zero missing values**. The column `_c39` had **1000 missing values**, indicating that it contained no usable information.

Action taken:

- `_c39` was dropped from the dataset.

2.2 Redundancy and Cardinality Checks

Unique value counts and ratios were calculated for all columns using:

- `value_counts()`
- `df.nunique() / len(df)`

Findings:

- `policy_number` and `incident_location` had **100% uniqueness**
- `insured_zip` and date columns had very high cardinality
- These variables act as identifiers rather than predictors

Action taken:

The following columns were dropped to reduce noise and overfitting risk:

- `policy_number`
- `incident_location`
- `insured_zip`

- `policy_bind_date`
- `incident_date`

These decisions were made **after inspection**, not arbitrarily.

3. Train–Validation Split

The cleaned dataset was split into:

- **70% training data**
- **30% validation data**

This split ensured:

- Model training, feature engineering, and resampling occurred **only on training data**
- Validation data remained untouched until final evaluation

This separation was maintained consistently throughout the notebook.

4. EDA on Training Data

EDA was performed **only on the training dataset** to avoid data leakage.

4.1 Univariate Analysis

Numerical Variables

Histograms were plotted for numerical columns such as:

- `total_claim_amount`
- `injury_claim`
- `vehicle_claim`
- `property_claim`
- `age`
- `months_as_customer`

Observations:

- Claim amount variables showed **right-skewed distributions**
- No extreme or invalid values were observed

- Age and customer tenure showed reasonable spreads

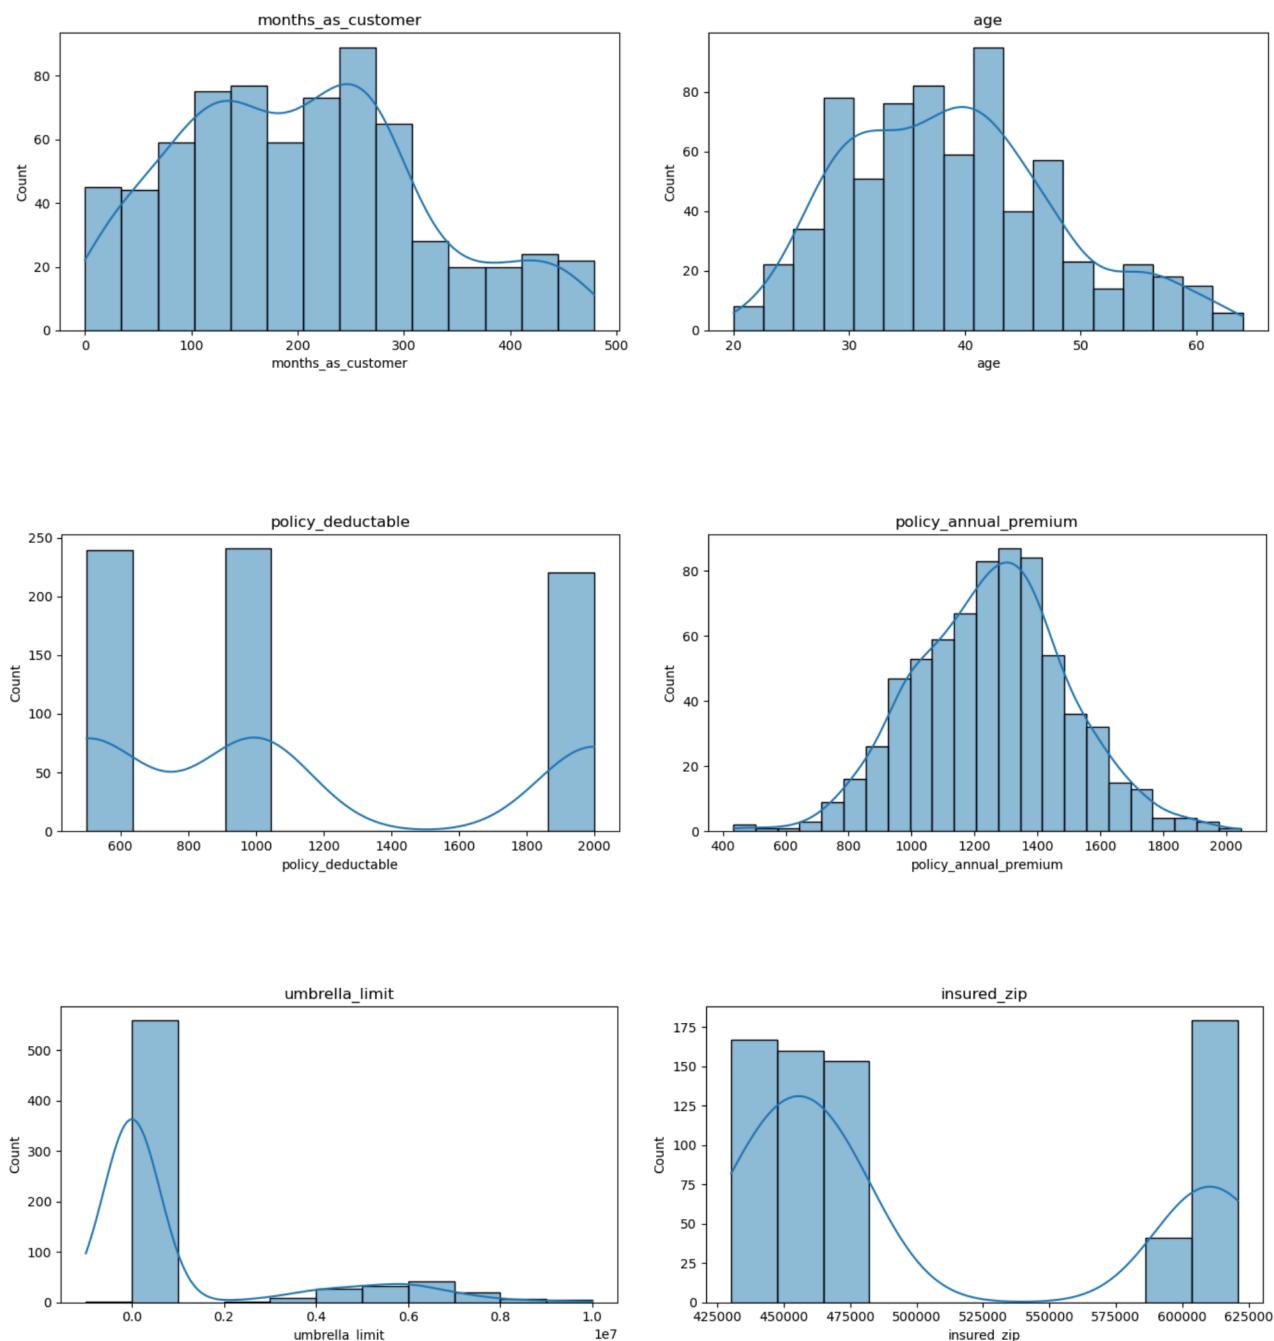
Categorical Variables

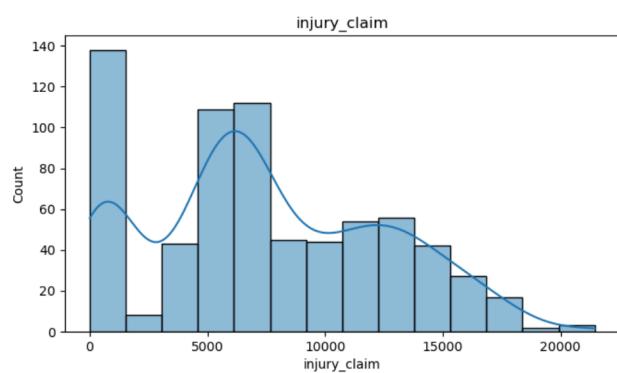
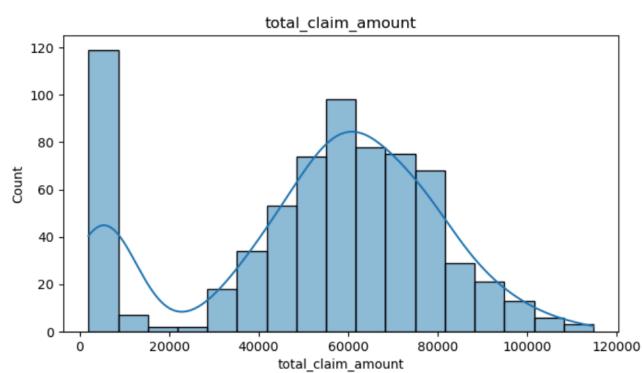
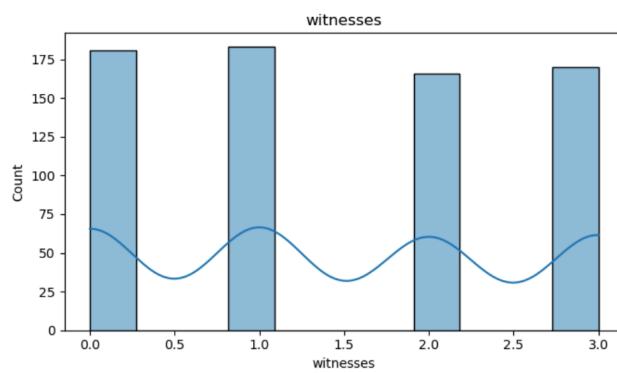
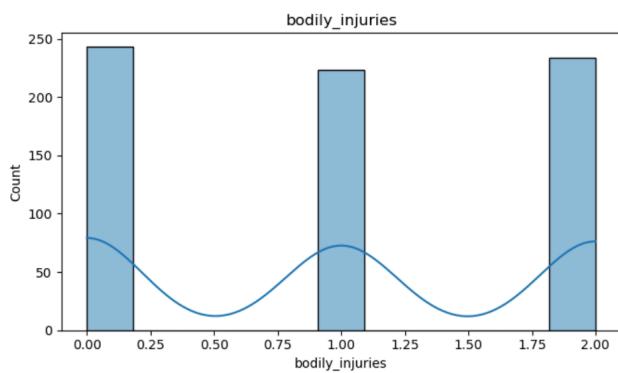
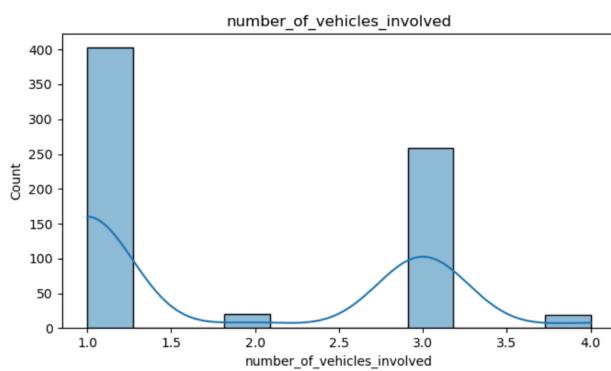
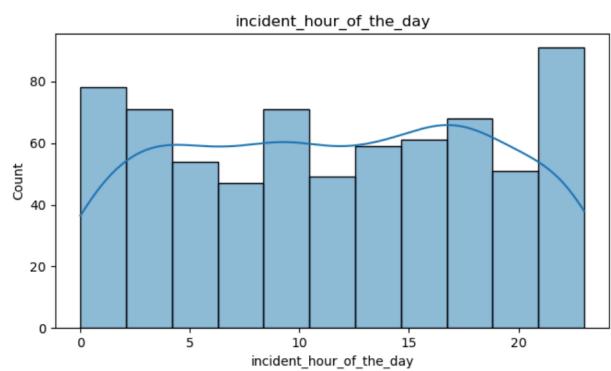
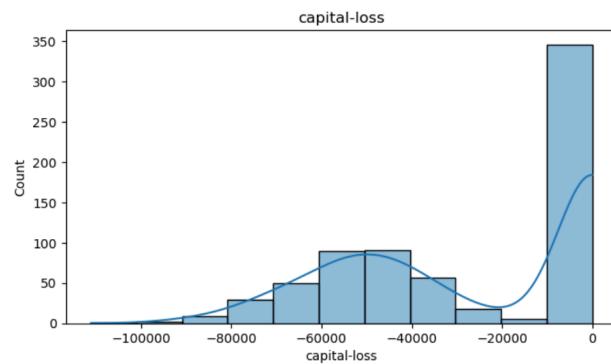
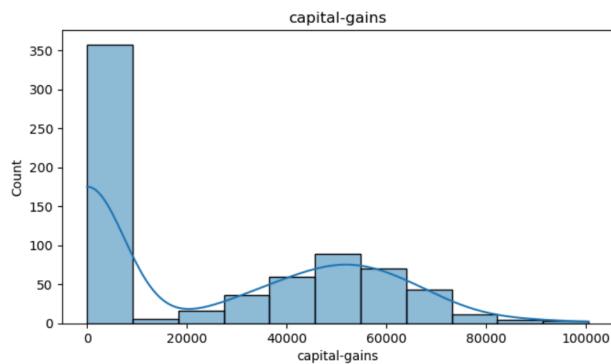
Count plots were generated for:

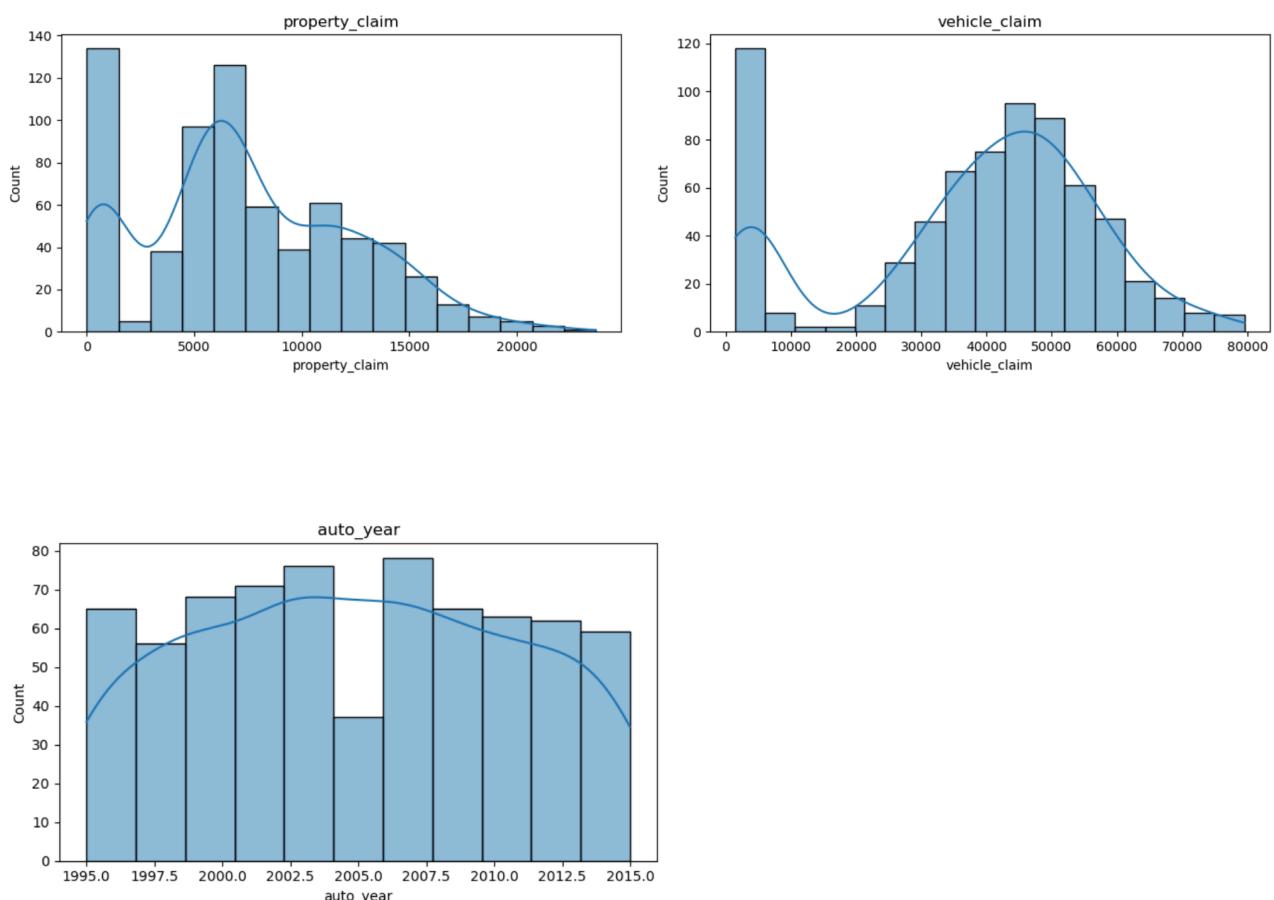
- `incident_type`
- `collision_type`
- `incident_severity`
- `insured_sex`
- `policy_state`

Several categories showed uneven distributions, indicating potential predictive signal.

Plots







4.2 Correlation Analysis

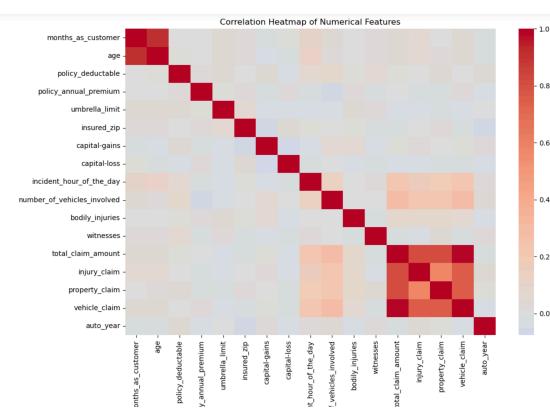
A correlation heatmap was generated for numerical variables.

Key observations:

- High correlation among claim amount components
- No severe multicollinearity across unrelated features

This supported proceeding with both logistic regression and tree-based models.

Plots



4.3 Class Balance Check

The target variable distribution was examined using `value_counts()`.

Observation:

- Fraudulent claims were significantly fewer than non-fraudulent claims

This confirmed **class imbalance**, justifying resampling during model training.



```
Out[33]: N    0.752857
          Y    0.247143
Name: fraud_reported, dtype: float64
```

4.4 Bivariate Analysis

Fraud likelihood was computed for categorical variables using group-by logic:

- Incident severity
- Incident type
- Collision type
- Authorities contacted
- Policy attributes

Key patterns observed:

- Major damage incidents had the highest fraud likelihood**
- Certain collision and incident types showed elevated fraud rates

- Claim severity was strongly associated with fraud

These insights directly informed feature engineering and model choice.

5. EDA on Validation Data

The same EDA steps were repeated on the validation dataset:

- Univariate distributions
- Correlation analysis
- Class balance check
- Bivariate fraud likelihood analysis

Fraud Likelihood Analysis (Output):-

Fraud Likelihood by policy_state
 policy_state
 OH 0.258475
 IN 0.244344
 IL 0.238683
 Name: fraud_reported, dtype: float64

Fraud likelihood by policy_csl:
 policy_csl
 100/300 0.276860
 250/500 0.234127
 500/1000 0.228155
 Name: fraud_reported, dtype: float64

Fraud likelihood by insured_sex:
 insured_sex
 FEMALE 0.248663
 MALE 0.245399
 Name: fraud_reported, dtype: float64

Fraud likelihood by insured_education_level:
 insured_education_level
 MD 0.277778
 PhD 0.261905
 JD 0.258065
 College 0.250000
 Associate 0.242991
 Masters 0.242718
 High School 0.205357
 Name: fraud_reported, dtype: float64

Fraud likelihood by insured_occupation:
 insured_occupation

```
transport-moving    0.400000
exec-managerial   0.333333
craft-repair      0.301887
farming-fishing   0.272727
sales             0.263158
armed-forces     0.260870
tech-support      0.250000
machine-op-inspect 0.235294
other-service     0.204545
protective-serv   0.204545
prof-specialty    0.203390
priv-house-serv   0.176471
adm-clerical      0.166667
handlers-cleaners 0.147059
Name: fraud_reported, dtype: float64
```

Fraud likelihood by insured_hobbies:

```
insured_hobbies
chess        0.878788
cross-fit    0.826087
polo         0.333333
yachting     0.323529
hiking        0.300000
board-games   0.250000
base-jumping  0.250000
paintball     0.244898
reading       0.234043
video-games   0.222222
skydiving     0.210526
sleeping      0.192308
bungie-jumping 0.179487
exercise      0.157895
movies        0.135135
camping       0.108108
dancing        0.108108
basketball    0.105263
kayaking      0.095238
golf          0.073171
Name: fraud_reported, dtype: float64
```

Fraud likelihood by insured_relationship:

```
insured_relationship
other-relative  0.290076
unmarried      0.280374
wife           0.276786
not-in-family   0.247863
husband        0.212389
own-child      0.175000
Name: fraud_reported, dtype: float64
```

Fraud likelihood by incident_type:

```
incident_type
```

Single Vehicle Collision 0.292419
Multi-vehicle Collision 0.272727
Parked Car 0.089552
Vehicle Theft 0.084746
Name: fraud_reported, dtype: float64

Fraud likelihood by collision_type:
collision_type
Rear Collision 0.301435
Front Collision 0.282486
Side Collision 0.260638
? 0.087302
Name: fraud_reported, dtype: float64

Fraud likelihood by incident_severity:
incident_severity
Major Damage 0.596939
Total Loss 0.148718
Minor Damage 0.093117
Trivial Damage 0.064516
Name: fraud_reported, dtype: float64

Fraud likelihood by authorities_contacted:
authorities_contacted
Ambulance 0.304965
Fire 0.297297
Other 0.279720
Police 0.204082
None 0.083333
Name: fraud_reported, dtype: float64

Fraud likelihood by incident_state:
incident_state
OH 0.470588
SC 0.313609
NC 0.305882
PA 0.260870
VA 0.259740
NY 0.204420
WV 0.155405
Name: fraud_reported, dtype: float64

Fraud likelihood by incident_city:
incident_city
Arlington 0.303922
Columbus 0.270270
Springfield 0.256637
Northbrook 0.231707
Riverwood 0.230769
Hillsdale 0.221154
Northbend 0.206186
Name: fraud_reported, dtype: float64

Fraud likelihood by property_damage:

property_damage

? 0.289157

YES 0.280193

NO 0.176230

Name: fraud_reported, dtype: float64

Fraud likelihood by police_report_available:

police_report_available

? 0.260331

NO 0.254098

YES 0.224299

Name: fraud_reported, dtype: float64

Fraud likelihood by auto_make:

auto_make

Audi 0.361702

Mercedes 0.333333

Ford 0.307692

BMW 0.280000

Chevrolet 0.272727

Toyota 0.255814

Honda 0.243243

Suburu 0.241935

Jeep 0.224490

Volkswagen 0.224490

Saab 0.210526

Dodge 0.200000

Nissan 0.188679

Accura 0.137255

Name: fraud_reported, dtype: float64

Fraud likelihood by auto_model:

auto_model

Silverado 0.500000

X6 0.454545

F150 0.444444

Civic 0.428571

A5 0.409091

Tahoe 0.400000

Maxima 0.352941

C300 0.333333

ML350 0.333333

E400 0.333333

A3 0.320000

Forrestor 0.320000

Grand Cherokee 0.315789

Highlander 0.312500

M5 0.300000

X5 0.294118

Impreza 0.272727

```

Camry      0.266667
Fusion     0.250000
93         0.235294
Neon       0.227273
Passat     0.227273
92x        0.227273
Jetta      0.222222
Escape      0.222222
Accord     0.200000
RAM         0.178571
MDX         0.172414
Wrangler   0.166667
95         0.166667
Corolla    0.166667
Ultima     0.157895
Legacy      0.153846
RSX         0.111111
3 Series   0.083333
TL          0.076923
CRV         0.076923
Pathfinder  0.058824
Malibu     0.000000
Name: fraud_reported, dtype: float64

```

Results closely mirrored training data patterns, confirming:

- The train-validation split was representative
- No distributional shift was introduced
- Higher fraud likelihood is observed for certain categories such as Single Vehicle Collision, Rear and Front Collisions, and incidents with Major Damage.
- Claims involving Ambulance or Fire authorities also show relatively higher fraud probabilities compared to cases where no authorities were contacted.
- Some insured attributes such as specific occupations, hobbies, and vehicle models exhibit noticeably higher fraud likelihood, while other categories show relatively uniform probabilities across levels.

6. Feature Engineering

6.1 Handling Class Imbalance

Random oversampling was applied **only to the training data** using `RandomOverSampler`.

Result:

- Fraud and non-fraud classes were perfectly balanced in training data
- Validation data was intentionally left unchanged

6.2 Derived Features

The following engineered features were created:

- `injury_claim_ratio`
- `vehicle_claim_ratio`
- `property_claim_ratio`
- `claim_severity_score`

These features captured **relative claim behavior**, improving signal quality beyond raw amounts.

6.3 Encoding

Categorical variables were converted into dummy variables using
`pd.get_dummies(drop_first=True)`.

Result:

- Feature count increased from ~35 to ~148 columns
- This is expected behavior for dummy variable creation
- Training and validation datasets were aligned to ensure consistent feature space

6.4 Feature Scaling

Numerical features were scaled using `StandardScaler`:

- Scaler fitted on training data
- Same scaler applied to validation data

This ensured fair model learning and evaluation.

7. Model Building

7.1 Logistic Regression

Training

- Logistic regression was trained on resampled, scaled training data
- Target variable encoded as binary (0/1)

Interpretation

- Coefficients were examined
- VIF was calculated to check multicollinearity
- Predicted probabilities were generated

7.2 Logistic Regression Evaluation

Training Performance

- Accuracy $\approx \mathbf{90.7\%}$
- Sensitivity $\approx \mathbf{92.6\%}$
- Specificity $\approx \mathbf{88.8\%}$
- F1 $\approx \mathbf{90.9\%}$

Cutoff Optimization

- ROC curve plotted
- Multiple cutoffs tested
- Optimal cutoff selected based on sensitivity-specificity tradeoff

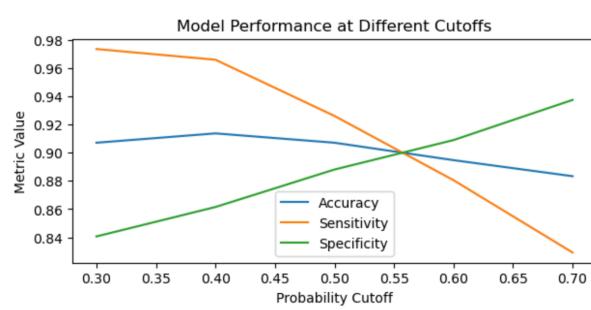
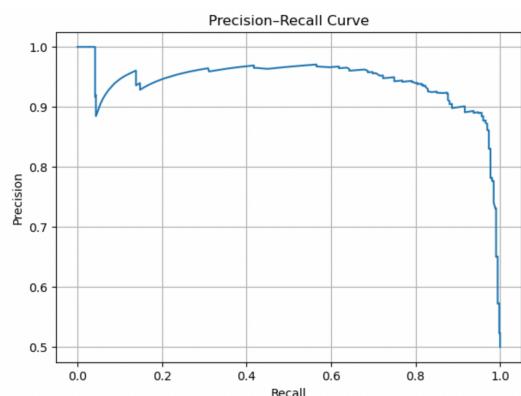
Validation Performance

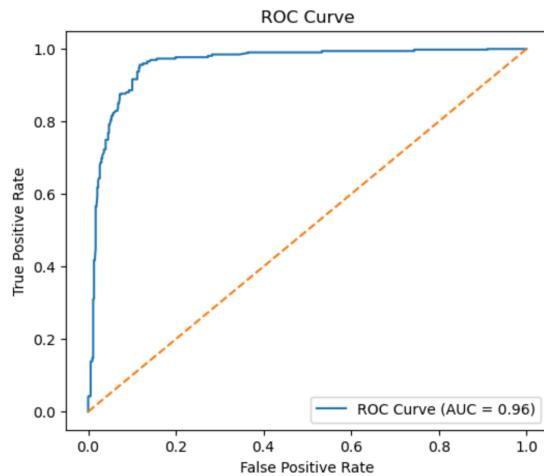
- Accuracy $\approx \mathbf{82.3\%}$
- Recall $\approx \mathbf{68.9\%}$
- F1 $\approx \mathbf{65.8\%}$

Conclusion:

Strong baseline, but recall on fraud cases could be improved.

Plots





7.3 Random Forest Model

Baseline Model

- Random Forest trained
- Feature importances extracted

Feature Selection

- Features with importance above mean retained

Cross-Validation

- 5-fold CV
- Mean accuracy $\approx 86.9\%$
- No overfitting observed

7.4 Hyperparameter Tuning

GridSearchCV was applied to tune:

- `n_estimators`
- `max_depth`
- `min_samples_split`

Best model retrained using optimal parameters.

8. Prediction & Model Evaluation

Logistic Regression (Validation)

- Accuracy $\approx 82\%$
- Lower fraud recall compared to training

Random Forest (Validation)

- Accuracy $\approx 84\%$
- Confusion matrix showed:
 - Higher fraud detection
 - Fewer false negatives

Random Forest consistently outperformed logistic regression on validation data.

9. Final Model Selection

Final Model Selection: Random Forest

Evidence-based justification for model selection:

- Higher validation accuracy
- Significantly better fraud recall
- Better F1-score
- Stable cross-validation performance

Training Performance (Random Forest)

- Accuracy: $\sim 88\%$
- Balanced confusion matrix with improved fraud detection

Validation Performance (Random Forest)

- Accuracy: $\sim 84\%$
- Confusion Matrix:
 - Correctly classified 193 non-fraudulent claims
 - Correctly identified 59 fraudulent claims

- 33 false positives
- 15 false negatives
-

Model Comparison and Selection - Logistic Regression vs Random Forest

Metric	Logistic Regression	Random Forest (Tuned)
Validation Accuracy	~82%	~84%
Recall (Fraud)	~69%	~80%
Precision	~63%	~64%
F1-score	~66%	~71%

Final Model Selected: Random Forest (Tuned)

Reasoning:

- Higher recall for fraudulent claims (business-critical)
- Better overall F1-score
- Improved generalisation on validation data
- Robust performance confirmed via cross-validation

10. Key Insights

- Fraud detection models must prioritise **recall over accuracy**, as missing fraudulent claims has higher cost.
- Probability cutoff optimisation significantly impacts model performance.
- Ensemble models like Random Forest capture non-linear relationships better than logistic regression.
- Feature importance helps improve interpretability and reduces overfitting.
- Hyperparameter tuning provides measurable performance gains.

11. Conclusion

This assignment successfully demonstrates a complete fraud detection pipeline, from data preparation to advanced model tuning and validation. While logistic regression provides a strong baseline, the tuned Random Forest model offers superior performance and better fraud detection capability.

The final selected model balances predictive accuracy with business relevance and is well-suited for real-world deployment in insurance fraud detection systems.